CLAUDE CERTIFICATION GUIDE · OCTOBER 2026 EDITION

CCAO-F

Claude Certified Associate - Foundations

The complete study guide, with 30 practice questions and an answer key

Full-length timed practice tests: preporato.com/certificates/claude-certified-associate-foundations

Updates, the practice test and the other three guides: github.com/preporato/claude-certification-guide

MIT license. Maintained by Preporato. Not affiliated with, endorsed by, or sponsored by Anthropic; Claude is a trademark of Anthropic, PBC. Exam facts reflect the program in October 2026; the Anthropic Partner Academy is authoritative.

Claude Certified Associate - Foundations (CCAO-F): The Complete Study Guide

A full preparation guide for Anthropic's non-technical certification. Part I teaches the material by topic. Part II re-indexes it by exam domain for revision. Part III covers trap answers, worked questions, and a final checklist.

Maintained by preporato.com. Facts checked against the official exam guide and Anthropic's program pages in October 2026.

Contents

- Introduction
- Exam format
- **Part I: The material**
  - Chapter 1: Evaluating outputs and catching hallucinations
  - Chapter 2: Fitting Claude into business workflows
  - Chapter 3: Governance and responsible use
  - Chapter 4: Prompting for business tasks
  - Chapter 5: Choosing models and features
  - Chapter 6: Projects, knowledge, and connectors
  - Chapter 7: Troubleshooting weak outputs
- **Part II: Exam domain notes**
  - All seven domains
- **Part III: Exam preparation**
  - The preparation ladder
  - The six trap answers
  - Exam-day strategy
  - Worked questions
  - 3-week study plan
  - Pre-exam checklist
  - Resources

Introduction

CCAO-F (Claude Certified Associate - Foundations) is the one Claude certification requiring no code, no API knowledge, and no engineering background. It validates that a professional can use the claude.ai product (the web and desktop application, as opposed to the developer platform) effectively and responsibly for real business work: drafting, research, analysis, planning, and process improvement.

Do not mistake non-technical for easy. The single largest domain is **evaluating outputs**, and more than a third of the exam sits in judgment territory: catching hallucinations before they reach a client, knowing when a human must own the final call, and recognizing which data never belongs in a chat. Candidates who prepare only on features routinely get caught by exactly those questions.

Who should take it: professionals in operations, marketing, project management, education, communications, and analysis who use Claude regularly and want a credential proving it. Anthropic recommends regular hands-on professional use, structured problem-solving, and a practical grasp of AI limitations.

Exam format

Questions	60, all scored
Duration	120 minutes to answer, about 135 minutes of seat time
Passing score	720 on a 100-1000 scaled score
Question style	Single-answer plus multiple-response (~a quarter of items), **no partial credit**
Delivery	Pearson VUE, test center or online proctored
Price	$99 per attempt, the most accessible entry in the program
Registration	Anthropic Partner Academy, open to people who work at Claude Partner Network organizations (sign in with a company email); passing earns a Credly badge
Languages	English or Japanese
Retakes	Full fee for every attempt; waits of 14, 30 and 90 days after the first, second and third failed attempts; at most 4 attempts in any 12 months
Validity	1 year; free non-proctored renewal before expiration, full proctored retake after a lapse
Scope	The claude.ai product only. Programming and API topics are deliberately excluded

Exam domains

Domain	Weight	~Questions	Core focus
1. Output Evaluation & Validation	21%	~13	Hallucination spotting, fact-checking, human review
2. Workflow Integration & Solution Design	16%	~10	Where Claude fits in real business processes
3. Governance, Risk & Responsible Use	15%	~9	Data sensitivity, policy, appropriate use
4. Prompting & Task Execution	14%	~8	Business prompts, decomposition, iteration
5. Product & Model Selection	12%	~7	Model tiers; chat vs Projects vs research
6. Configuration & Knowledge Management	12%	~7	Project instructions, knowledge, connectors
7. Troubleshooting & Optimization	10%	~6	Diagnosing weak outputs, systematic fixes

Part I: The material

Chapter 1: Evaluating outputs and catching hallucinations

The heaviest domain, and one habit underneath it: **every output is a draft that must earn trust before it leaves your screen.**

1.1 What a hallucination looks like

A hallucination is a confident, plausible statement that is factually wrong or unsupported by any source. The exam expects recognition from symptoms alone:

- **Highly specific numbers, dates, or statistics with no source.** Precision signals confidence and reveals nothing about accuracy.
- **Named studies, publications, or quotes you cannot locate.** Fabricated citations read exactly like real ones until checked.
- **Details beyond the material you provided.** If the summary of last quarter's report mentions figures the report never contains, Claude filled the gap itself.
- **Uniform confidence across every claim.** The weakest guesses arrive in the same fluent tone as established facts, so tone can never substitute for verification.
- **Answers about events after the model's knowledge cutoff** with no research mode or attached source grounding them.

A worked example, in the form the exam uses. You attached this excerpt from a quarterly report:

> Q2 revenue was $4.2M, up from Q1. The Enterprise segment led growth. We closed 14 new enterprise accounts during the quarter.

Claude's summary reads:

> Revenue grew 18% quarter-over-quarter to $4.2M, driven by the Enterprise segment, which closed 14 new accounts including two Fortune 500 logos. According to a Gartner study, this growth rate places the company in the top decile of its peer group.

Three hallucinations, each matching a red flag: **"18%"** is a specific figure the source never states (the source says only "up from Q1"), **"two Fortune 500 logos"** goes beyond the provided material, and **the Gartner study** is a citation you will not locate. The revenue figure, the segment, and the account count are grounded. The tone is identical across all six claims, which is exactly why tone cannot be the filter.

1.2 The validation sequence

For any output that carries consequences: verify every name, number, date, and quotation against the original source. Ask Claude directly which part of the attached document supports a specific claim; a grounded answer can point to its evidence, and a hallucinated one cannot. Cross-check external facts

in an independent reference before repeating them. And watch for bias as well as error: one-sided analysis, stereotyped examples, and missing counterarguments all count as evaluation failures.

1.3 When human review is mandatory

Anything external-facing (client emails, published copy, press materials), anything with legal, financial, or health consequences, anything feeding a decision about people (hiring, performance, vendor selection), and anything in a regulated workflow. The exam rewards keeping a person accountable at those points regardless of how good the output looks.

1.4 Choosing an output format

Inline chat response for quick answers and short summaries. An **Artifact** (a separate panel holding a standalone deliverable) for anything you will revise and share. A structured format (table, bulleted list) for output feeding another step, because structure makes checking easier.

Chapter 2: Fitting Claude into business workflows

Design questions start from the task: state the deliverable, the audience, the source material, and what success looks like before opening a chat. Jumping straight to prompting without defining the task is wrong at this level.

Strong fit	Weak fit
First drafts of emails, briefs, reports, plans	Final authority on factual accuracy
Summarizing meetings, long documents, research	Decisions requiring human accountability (hiring, legal, spend)
Reformatting and repurposing content across channels	Precise high-volume calculations a spreadsheet should own
Brainstorming options, angles, counterarguments	Current events and live data, unless research mode grounds them

Three principles carry the domain:

- **Keep the person at the judgment step.** Claude produces drafts and analysis; a human makes the decision and signs off on facts.
- **Communicate value and limits together.** Presenting Claude to stakeholders pairs a concrete outcome ("hours saved weekly on meeting summaries") with an honest limitation ("outputs are probabilistic and require review").
- **Pilot before rollout.** One workflow, measured before and after, then expand. Big-bang adoption answers lose to measured pilots.

Chapter 3: Governance and responsible use

One habit: **check the policy before you paste.** When unsure whether data is sensitive, treat it as sensitive.

Area	Do	Don't
Data sensitivity	Check your organization's AI policy before sharing internal data	Paste customer records, financials, or trade secrets for convenience
Personal data	Anonymize or summarize before sharing when the task allows	Include names, contacts, or health information you do not need
Appropriate use	Draft, summarize, analyze, brainstorm	Let Claude be sole decision-maker on hiring, legal, or medical questions
Regulated content	Route legal/financial/compliance outputs through qualified reviewers	Publish regulated advice on the strength of a fluent answer
Transparency	Follow your organization's AI-disclosure rules	Present AI-generated work as independently verified when it is not
Policy conflicts	Escalate ambiguous use cases to the policy owner	Assume approval because a tool is available

Two decision rules cover most scenarios: if an output affects people's **rights, money, health, or legal standing**, a qualified human owns the final call. If data would cause harm were it **leaked**, the policy check comes before any productivity benefit.

Chapter 4: Prompting for business tasks

4.1 The five-part business prompt

Role or context ("You are helping a project manager prepare a stakeholder update"), the task, the source material, the constraints (audience, length, tone), and the output format. "Summarize these meeting notes for the executive team in five bullet points, each naming an owner and a deadline" beats "summarize this" on every axis the exam measures.

Before and after, on the same task:

> **Weak:** "Summarize these meeting notes."

> **Strong:** "You are helping a project manager prepare a stakeholder update. Summarize the attached meeting notes for the executive team in five bullet points. Each bullet names the decision made, the owner, and the deadline. Formal tone, no jargon, nothing that was only discussed without a decision."

The strong version carries all five parts: role, task, source, constraints, format. Every improvement in the output is traceable to one of those five additions, which is how the exam expects you to diagnose weak prompts in reverse.

4.2 Task decomposition

Split a large deliverable into a sequence of smaller prompts, each producing something checkable before moving on. A quarterly report becomes: agree on an outline, draft one section at a time from its source material, then a final pass for tone and consistency. Decomposition is the tested answer whenever a scenario involves a multi-part deliverable, a long document, or a shallow result from one giant prompt.

4.3 Iteration

The first response is a starting point. Steer with specific feedback: "cut this to half the length, make the tone more direct, keep the three statistics" gives Claude something to act on. Regenerating and hoping does not, and the exam consistently rewards targeted refinement over rerolls.

4.4 Prompt strategy by task type

Task type	Adjust the prompt to
Analysis	Provide sources, name the criteria to weigh, ask for reasoning before conclusions
Research	Use research mode, require checkable sources
Drafting	Specify audience, tone, length, format; share an example of good
Brainstorming	Ask for volume and variety first, evaluate and narrow second

Chapter 5: Choosing models and features

5.1 Model selection

Model	Optimizes for	Pick it for
Haiku	Speed and low cost	Quick, well-defined tasks at volume: reformatting, short summaries, first-pass triage
Sonnet	Balance of quality, speed, cost	The everyday default: drafting, analyzing reports, synthesizing notes
Opus	Deepest reasoning, highest quality	Complex high-stakes work: multi-document strategy, plans where an error is expensive

Default to Sonnet and move on evidence: up to Opus when the task genuinely demands deeper reasoning, down to Haiku when the task is simple and speed beats polish. The most powerful model for every task is a recurring wrong answer.

5.2 Feature selection

Feature	What it is	Reach for it when
Chat	A single standalone conversation	One-off question or quick draft, no saved setup
Projects	A workspace with standing instructions and uploaded knowledge shared across its chats	The same context, documents, or standards recur across sessions

Feature	What it is	Reach for it when
Research mode	Claude searches current sources and cites findings	You need up-to-date, verifiable information
Artifacts	A side panel holding a standalone, editable deliverable	Producing a document or table to refine over rounds and share

5.3 Context limits and the three moves

The context window is the amount of conversation and attached material Claude can consider at once, and long sessions strain it. Three moves cover the exam's scenarios: **restart** a fresh chat when a long conversation drifts or ignores instructions; **summarize** the decisions so far and carry the summary into the new chat; **persist** durable reference material into Project knowledge so no single conversation holds it.

Chapter 6: Projects, knowledge, and connectors

A Project has two configurable layers, and the classic exam question is knowing which layer a piece of setup belongs in:

- **Project instructions**: standing directions applied to every chat in the workspace (role, tone standards, required formats).
- **Project knowledge**: uploaded documents (style guides, product sheets, past reports) Claude draws on in every conversation.

You want to...	Put it in
Apply a tone, role, or format standard to every chat	Project instructions
Make reference documents available to every chat	Project knowledge
Give details that matter only for today's task	The individual chat prompt
Pull from email or files that change frequently	A connector (Gmail, Google Drive)

What good Project instructions look like, for a marketing team's content workspace:

> You are a content editor for Meridian Software's marketing team. In every chat:
>
> - Follow the brand voice guide in Project knowledge; flag any request that conflicts with it
> - Audience is IT decision-makers; avoid consumer-marketing tone
> - All statistics must come from the attached fact sheet; if a claim is not there, say so instead of estimating
> - Default deliverable format: headline options first, then body copy, then a three-line summary for approvals

Standing role, standards, a guardrail against invented statistics, and a default format. The documents those instructions reference (the voice guide, the fact sheet) live in Project knowledge; today's campaign brief goes in the individual chat.

Connectors link Claude to live sources with your permission; reach for one when material lives in those tools and changes often enough that a static upload would go stale. Configuration needs maintenance: revisit instructions when outputs drift from standards, and update or remove knowledge files when the underlying documents change, because stale knowledge produces confidently outdated answers.

Chapter 7: Troubleshooting weak outputs

The exam gives a symptom and asks for the most effective fix:

Symptom	Likely cause	First fix
Generic output that could apply to any company	Prompt lacks context and constraints	Add role, audience, sources, and an example of good
Instructions ignored late in a long chat	Context overload and drift	Restart with a summary; move standing rules into Project instructions
Confident but wrong claims about your business	No source material provided	Attach documents or connect the source, ask Claude to cite them
Wrong format every time	Format described vaguely	Show a concrete example of the structure
Quality varies across teammates on the same task	No shared setup	A shared Project with instructions and knowledge
Half of a multi-part request ignored	Too many asks in one prompt	Decompose and run steps in sequence

The optimization rule tying the domain together: **change one variable at a time.** Adjust the prompt, or the model, or the knowledge, compare, then change the next thing. Changing several at once tells you nothing about what worked.

Part II: Exam domain notes

Domain 1: Output Evaluation & Validation (21%)

Know: the five hallucination red flags; the validation sequence; the mandatory-human-review list; which output format fits which use.

Be able to: spot a fabricated citation or ungrounded figure from symptoms; ask Claude to point at its supporting evidence; name the review gate a scenario requires.

Anti-patterns: trusting fluency; skipping verification because the output "looks right"; treating tone as a signal of accuracy.

Domain 2: Workflow Integration & Solution Design (16%)

Know: the strong-fit/weak-fit table; the human-at-the-judgment-step division; value-plus-limits stakeholder framing; pilot-then-expand.

Be able to: define deliverable, audience, sources, and success before prompting; place Claude at the draft stage of a real process.

Anti-patterns: Claude as final factual authority; big-bang rollouts; benefits pitched without limits.

Domain 3: Governance, Risk & Responsible Use (15%)

Know: the do/don't matrix; the two decision rules (rights/money/health/legal → human owns the call; leak-harmful data → policy before productivity).

Be able to: classify data sensitivity conservatively; escalate ambiguous cases to the policy owner; route regulated outputs through qualified reviewers.

Anti-patterns: paste-first; sole-decision-maker Claude; presenting AI work as independently verified.

Domain 4: Prompting & Task Execution (14%)

Know: the five-part prompt; decomposition triggers; iteration with targeted feedback; the per-task-type adjustments.

Be able to: rewrite "summarize this" into a complete prompt; sequence a large deliverable into checkable steps.

Anti-patterns: one-giant-prompt; regenerate-and-hope.

Domain 5: Product & Model Selection (12%)

Know: the three model rules and the Sonnet default; the four feature triggers; the three context moves (restart, summarize, persist).

Be able to: match a business task to a model on speed/cost/quality; pick chat vs Project vs research mode vs Artifact from scenario wording.

Anti-patterns: strongest-model-everywhere; rebuilding context by hand every session.

Domain 6: Configuration & Knowledge Management (12%)

Know: instructions vs knowledge vs chat prompt vs connector; connector triggers; configuration maintenance.

Be able to: place any setup detail in its correct layer; recognize stale knowledge as a cause of confidently outdated answers.

Anti-patterns: standing rules retyped per chat; static uploads of fast-changing sources.

Domain 7: Troubleshooting & Optimization (10%)

Know: all six symptom-cause-fix rows; one-variable-at-a-time.

Be able to: pick the first fix from the symptom, and diagnose before changing anything.

Anti-patterns: multi-variable changes; fixing generic output by regenerating.

Part III: Exam preparation

The preparation ladder

Four layers, in order. Each builds on the previous, and the free layers take you surprisingly far. Anthropic retired its official practice exam when the exams moved to Pearson VUE; the sample questions in the official exam guide are what remains.

Step	Resource	Cost
1. Learn the vocabulary	Anthropic Partner Academy courses, the official exam guide, and the Anthropic Academy courses mapped in Resources	Free
2. Go to the source	The claude.ai Help Center, the primary source for this exam	Free
3. First calibration	30 practice questions with explanations in this repo, also as a timed practice test	Free
4. Full calibration	Six full-length 60-question practice tests and a 500-card flashcard deck on Preporato, built on the exact domain weights; the page has free sample questions to try first	Paid

Steps 1-3 cost nothing and tell you honestly where you stand. Step 4 is how you close the gap they reveal.

The six trap answers

Trap	Why it fails
Trust-the-output	Fluency is not verification; consequential output gets validated
Strongest-model-everywhere	Ignores speed and cost; choice matches complexity and stakes
One-giant-prompt	Multi-part deliverables get decomposed into checkable steps
Paste-first	The policy check precedes the productivity win
Regenerate-and-hope	Targeted feedback beats rerolls
Rebuild-every-session	Recurring context belongs in a Project

On multiple-response items, apply the traps per option: one paste-first or trust-the-output pick among your selections costs the whole question, because there is no partial credit.

Exam-day strategy

120 minutes for 60 questions is two minutes each, and CCAO-F stems are short, so time pressure is mild. The risk is pattern-matching on features when the question is really about judgment: before answering, ask which domain this is, and if the scenario involves facts, people, or sensitive data, expect the answer to include a safeguard. Flag long deliberations, return at the end, change answers only on certain error.

Worked questions

Ten fresh questions in the exam's style. (Written for this guide; timed full-length practice is linked in Resources.)

Q1. Claude summarizes your attached quarterly report and includes a revenue growth percentage. The report contains no such percentage. What happened, and what is the correct response?

- A) Claude calculated it from the report's figures; include it
- B) A hallucination filled a gap in the source; remove or verify the figure against the original before the summary goes anywhere
- C) The model used more recent data; keep the number
- D) A formatting error; regenerate the summary

Answer: B. Details beyond the provided material are a named red flag. A assumes a calculation nobody verified, C invents a data source chat Claude does not have, and D rerolls without addressing the accuracy problem.

Q2. A manager wants Claude to rank the team's five sales reps for a promotion decision, using performance notes. The correct guidance?

- A) Fine, if the notes are anonymized first
- B) Fine, if Opus is used for deeper reasoning
- C) Claude can summarize and organize the notes, but a decision about people requires the manager to own the judgment and the final call
- D) Refuse to use Claude anywhere near performance data

Answer: C. Decisions about people (hiring, performance) sit on the mandatory-human-judgment list. A solves a privacy slice of the problem and misses the accountability core, B upgrades the model without changing who decides, and D overcorrects past the do-column uses that remain fine.

Q3. A colleague pastes a customer list with emails and contract values into a chat "to save time formatting it." What should have happened first?

- A) Switching to a more capable model for accuracy
- B) Checking the organization's AI data policy, and anonymizing or summarizing the sensitive fields if the task allows
- C) Turning on research mode for verification
- D) Nothing; formatting is a low-risk task

Answer: B. Paste-first is the trap: leak-harmful data (customer records, financials) triggers the policy check before any convenience. The task being "just formatting" changes nothing about what was shared.

Q4. Your prompt "write a blog post about our new product" returned something generic enough to describe any company's product. The most effective fix?

- A) Regenerate until a better version appears
- B) Switch to the most capable model

- C) Rewrite the prompt with role, audience, tone, length, the product sheet attached, and an example post you like
- D) Split the request across three chats

Answer: C. Generic output means the prompt lacked context and constraints; the five-part prompt is the fix. A is regenerate-and-hope, B is strongest-model-everywhere, and D decomposes without adding the missing context.

Q5. Every week you produce the same client status update: same tone rules, same format, drawing on the same background documents. Today you retype the instructions and re-upload the files. The better setup?

- A) A Project with the tone and format in instructions and the background documents in knowledge, so every chat starts configured
- B) A saved text file of prompts to paste each time
- C) One very long chat reused forever
- D) Research mode, which remembers sources

Answer: A. Recurring context is the Projects trigger, and the instructions/knowledge split matches standing rules vs reference material. B still rebuilds by hand, C walks into context drift, and D misdescribes research mode.

Q6. (Select TWO) Which two situations make human review mandatory rather than optional?

- A) A draft agenda for an internal team meeting
- B) A response to a customer complaint about billing, sent under the company's name
- C) A brainstorm of taglines for an internal workshop
- D) A summary feeding a decision on which vendor contract to terminate
- E) A reformatted version of your own notes

Answer: B and D. External-facing communication with financial context (B) and input to a consequential decision about money and vendors (D) both sit on the mandatory-review list. A, C, and E are low-stakes internal drafts.

Q7. Twenty exchanges into a long chat, Claude starts ignoring the formatting rules you set at the beginning. Best response?

- A) Repeat the rules more forcefully in the next message
- B) Restart a fresh chat carrying a summary of decisions so far, and move the standing rules into Project instructions so they persist
- C) Switch to a faster model
- D) Ask Claude why it is ignoring you

Answer: B. Late-chat instruction loss is context drift; the restart-summarize-persist trio is the fix, and instructions belong in the layer that survives across chats. A treats a structural problem as an emphasis problem, and C changes speed rather than context.

Q8. You need current pricing for three competitors, announced within the last month, for a comparison table you will present. Which approach fits?

- A) Ask Claude directly and trust its training knowledge
- B) Use research mode so Claude searches current sources with citations, then verify the cited pages before presenting
- C) Use Opus, whose knowledge is more recent
- D) Ask three times and keep the consistent answers

Answer: B. Recent, verifiable information is the research mode trigger, and the citations still get checked before an external-facing deliverable. A and C rely on training data for post-cutoff facts (model capability does not change the cutoff), and D mistakes consistency for accuracy.

Q9. Your team's five analysts produce weekly client reports with Claude, and quality varies wildly: different tones, different structures, statistics of unclear origin. Which single change addresses all three symptoms?

- A) Have the strongest analyst review everyone's output before it ships
- B) A shared Project: tone and structure standards in the instructions, the approved data sources in knowledge, used by all five analysts
- C) Standardize on the most capable model
- D) A shared document of prompt examples analysts can copy from

Answer: B. Variance across teammates on the same task is the no-shared-setup symptom, and a shared Project fixes tone (instructions), structure (instructions), and data provenance (knowledge) in one move. A adds a bottleneck without fixing the source, C changes quality but not consistency, and D depends on five people copying correctly forever.

Q10. A colleague asks Claude for your company's Q3 market share, gets "23.4%, per IDC's industry tracker," and pastes it into a board slide. No source was attached to the chat. What should have happened?

- A) Nothing; the citation makes it credible
- B) Ask Claude to double-check its own number in the same chat
- C) Treat it as unverified: a specific figure with a named study and no attached source is the hallucination profile, so verify against the actual IDC data (or use research mode with checkable citations) before it reaches the board
- D) Rephrase the question and compare the two answers

Answer: C. Specific number, named study, no grounding: two red flags at once, headed for an external-facing, high-stakes surface, which also makes human verification mandatory. A trusts the tell that most needs checking, B asks the same ungrounded process to grade itself, and D mistakes consistency for accuracy.

3-week study plan

- **Week 1: official baseline.** Check that your employer is in the Claude Partner Network (registration needs your company email), complete the Partner Academy associate courses, read the exam guide, and map yourself against the seven domains.
- **Week 2: deliberate product practice.** Work through the claude.ai Help Center systematically, and use every feature the blueprint names on your own real work: build a Project with instructions and knowledge, produce Artifacts, run research mode, connect a live source, and compare the three models on one task. Practice the judgment side deliberately: run the validation sequence on a real output, and apply the governance matrix to your own team's data.
- **Week 3: calibrate.** Start with the 30 practice questions in this repo to see the format, then take full-length timed practice tests, with per-domain review and a final pass on the two judgment domains (evaluation and governance), which catch feature-focused candidates. The night before: Part III here plus the cheat sheet.

Expanded version: 3-week CCAO-F study plan.

Pre-exam checklist

- ☐ I can name five hallucination red flags and what each signals
- ☐ I can run the validation sequence and say when human review is mandatory
- ☐ I know which output format fits a quick answer, a shared deliverable, and a downstream process
- ☐ I can name Claude's strong fits and weak fits in a business workflow
- ☐ I can present Claude to stakeholders with a concrete benefit and an honest limitation
- ☐ I can apply the governance do/don't matrix to a data-sensitivity scenario
- ☐ I know the two governance decision rules (rights/money/health/legal; leak-harm)
- ☐ I can write a five-part business prompt
- ☐ I can decompose a large deliverable into checkable steps and iterate with specific feedback
- ☐ I can choose Haiku, Sonnet, or Opus from speed, cost, and quality needs
- ☐ I can pick chat, Project, research mode, or Artifact from a scenario trigger
- ☐ I know when to restart, summarize, or persist as a chat grows long
- ☐ I can place any setup detail into instructions, knowledge, a prompt, or a connector
- ☐ I can match all six troubleshooting symptoms to their first fix
- ☐ I change one variable at a time when improving a workflow
- ☐ I have the six traps down cold and apply them per option on multiple-response items
- ☐ I have taken at least two full-length timed practice tests and reviewed every miss

Resources

Official (free):

- Anthropic Partner Academy: the exam guide, prep courses, and registration
- claude.ai Help Center: the primary source for this exam
- Anthropic Academy courses, mapped to this exam's domains:
  - Claude 101: the product tour this exam assumes

- AI Fluency: Framework & Foundations: delegation, description, discernment, and diligence; a direct match for the output evaluation and governance domains, with a certificate of completion
- AI Capabilities and Limitations: the judgment layer behind hallucination spotting and appropriate use
- Audience variants exist too (small businesses, educators, students) if one matches your role

Free articles (preporato.com):

- CCAO-F complete guide · Cheat sheet · Domains breakdown · How to pass first attempt · CCAO-F vs CCDV-F

Practice:

- 30 practice questions with explanations in this repo, also as a timed practice test (free)
- The CCAO-F page on Preporato: free sample questions to start, then the full tests below
- (Paid) Six full-length 60-question practice tests and a 500-card flashcard deck built on the exact domain weights

---

Not affiliated with Anthropic. Exam details reflect the program as of October 2026; the Partner Academy is authoritative. Found an error? Open an issue.

Practice questions

30 questions for CCAO-F, written for this guide in the exam's style and spread across its domains by the official weights. Questions that ask for more than one answer say how many. The answers and explanations follow in the next section.

The same questions run as a timed test with a score for every domain: practice test.

Question 1 *(Output Evaluation and Validation)*

Claude summarized a 30-page vendor contract you attached and wrote: 'The contract auto-renews for 24 months unless cancelled 90 days before expiry.' This line is going into an email to your director. What is the most efficient way to check it?

- **A.** Ask Claude to quote the exact clause and give its section number, then read that clause in the contract yourself
- **B.** Ask Claude whether it's sure
- **C.** Regenerate the summary and see whether the line comes back the same
- **D.** Send it as is, because the contract was attached

Question 2 *(Output Evaluation and Validation)*

Claude drafted a job ad for a software engineer that uses phrases like 'young, energetic team' and 'digital native.' What is the issue, and what should you do?

- **A.** Nothing; these phrases are standard in job ads
- **B.** The wording signals an age preference, which can discourage qualified applicants and create legal risk; ask for inclusive wording and check it against your hiring policy
- **C.** The ad is too long; ask for a shorter version
- **D.** The tone is too casual for a job ad; ask for a formal rewrite

Question 3 *(Output Evaluation and Validation)*

Claude built a pricing comparison table with totals calculated from 40 line items you pasted in. The table is going into a client proposal. What validation is needed?

- **A.** Round the totals so small differences don't matter
- **B.** Nothing, because you supplied the numbers
- **C.** Recompute the totals independently, for example in a spreadsheet, before the proposal goes out
- **D.** Ask Claude to double-check its math in the same chat

Question 4 *(Output Evaluation and Validation)*

Claude wrote a market summary with no sources attached. Which TWO details are hallucination red flags? (Select TWO)

- **A.** 'According to a 2025 McKinsey survey, 73% of retailers plan to...' with no link or attached source
- **B.** A general explanation of what market share means
- **C.** A precise figure for next quarter's market size, stated as fact

- **D.** A note that Claude can't see your internal sales data
- **E.** A suggestion to verify the figures before publishing

Question 5 *(Output Evaluation and Validation)*

Claude summarized a customer interview and wrote: 'The customer is frustrated with onboarding and plans to cancel.' The transcript mentions onboarding friction but says nothing about cancelling. What is the problem?

- **A.** Nothing; it's a reasonable reading of the frustration
- **B.** The summary is too short
- **C.** The summary needs a more formal tone
- **D.** An unsupported inference is presented as fact; edit the line to what the transcript supports, or label it clearly as an interpretation

Question 6 *(Output Evaluation and Validation)*

Research mode returns a summary of a new data regulation with eight cited sources. Your compliance team will base a plan on it. Which TWO steps are needed before anyone uses it? (Select TWO)

- **A.** Ask Claude to rate the reliability of each source
- **B.** Have the compliance owner review the conclusions before the plan is built on them
- **C.** Open the cited sources and confirm that each key claim matches what the source says, and that the sources are current and authoritative
- **D.** Run research mode again and keep the claims that appear twice
- **E.** Nothing; the citations prove the summary is accurate

Question 7 *(Workflow Integration and Solution Design)*

Your manager wants to roll Claude out to all 200 employees next week for every kind of task. What do you recommend?

- **A.** Let each employee figure it out without guidance
- **B.** Roll it out to everyone at once to build momentum
- **C.** Pilot one or two well-defined workflows with a small group, measure before and after (time spent, error rate), then expand
- **D.** Wait until AI tools stop making mistakes

Question 8 *(Workflow Integration and Solution Design)*

Which TWO tasks are strong fits for Claude in a marketing team's workflow? (Select TWO)

- **A.** Turning a webinar transcript into a blog draft and three social posts
- **B.** Deciding which agency gets a $200,000 contract
- **C.** Reconciling 5,000 ad-spend transactions with no independent check
- **D.** Reporting a competitor's announcement from yesterday, from memory
- **E.** Drafting five subject-line variants for an A/B test

Question 9 *(Workflow Integration and Solution Design)*

You're presenting a Claude pilot to leadership. Which framing is most appropriate?

- **A.** 'Meeting summaries now take 10 minutes instead of 45. The meeting owner still reviews each one, because outputs can contain errors.'
- **B.** 'Claude eliminates errors from our summaries.'
- **C.** 'Claude will replace the operations team.'
- **D.** 'The benefits can't really be measured.'

Question 10 *(Workflow Integration and Solution Design)*

A colleague opens Claude and types 'help me with the Q3 board deck.' The results are unfocused. What should be defined before prompting?

- **A.** The slide color scheme
- **B.** Nothing; Claude should infer it
- **C.** Which model to use
- **D.** The deliverable, the audience, the source material and what a good result looks like

Question 11 *(Workflow Integration and Solution Design)*

A legal team handles NDAs in five steps: intake, first-pass review against the team's playbook, negotiation, approval and signature. Where does Claude add the most value with the right safeguards?

- **A.** Signing on the company's behalf
- **B.** First-pass review: compare each NDA to the playbook and flag deviations for a lawyer to decide on
- **C.** Final approval
- **D.** Negotiating directly with the other party, unsupervised

Question 12 *(Governance, Risk, and Responsible Use)*

You want Claude to summarize an internal incident report that contains customers' names and email addresses. Your company's AI policy allows internal documents and prohibits customer personal data. What do you do?

- **A.** Paste it in and ask Claude to ignore the names
- **B.** Paste it in, because it's an internal document
- **C.** Use the most capable model so the data stays safe
- **D.** Remove or anonymize the customer details first (the summary doesn't need them), then summarize

Question 13 *(Governance, Risk, and Responsible Use)*

A sales rep wants to upload a prospect's RFP, marked confidential, so Claude can help draft the response. Your AI policy says nothing about third-party confidential documents. What is the best action?

- **A.** Ask the policy owner or legal before uploading, and treat the document as sensitive until there's an answer

- **B.** Upload only the pages that look less sensitive
- **C.** Upload it, since the policy doesn't forbid it
- **D.** Rewrite the RFP from memory and upload that

Question 14 *(Governance, Risk, and Responsible Use)*

Your organization requires disclosure when AI helps produce external content. You drafted a customer newsletter with Claude and then edited it heavily. What applies?

- **A.** Disclose only if a customer asks
- **B.** Follow the disclosure rule as written, because your edits don't change how the draft was produced
- **C.** Describe the newsletter as written and fact-checked without AI
- **D.** Skip disclosure, because your edits make it your own work

Question 15 *(Governance, Risk, and Responsible Use)*

Which TWO uses make Claude the decision-maker in a situation where a qualified person must own the call? (Select TWO)

- **A.** Brainstorming names for a team offsite
- **B.** Using Claude's ranking alone to decide which employees are laid off
- **C.** Drafting three versions of an internal FAQ for a manager to choose from
- **D.** Sending Claude's answer about a customer's legal rights to the customer without legal review
- **E.** Summarizing your own meeting notes

Question 16 *(Governance, Risk, and Responsible Use)*

A financial advisory firm uses Claude to draft client emails that explain investment options. Which control is essential?

- **A.** Using the most capable model
- **B.** A disclaimer at the bottom of every email, with no review
- **C.** Compliance review of client-facing investment content before it's sent, including a check of every figure
- **D.** Sending the drafts faster to improve service

Question 17 *(Prompting and Task Execution)*

Your prompt 'Write an email about the new expense tool' produced something generic. Which TWO additions would improve it most? (Select TWO)

- **A.** The tool's key facts, attached or pasted in: launch date, how to get access, where to get help
- **B.** A switch to the most capable model
- **C.** A request for ten versions
- **D.** Who the email is for and what it should get them to do
- **E.** 'Make it amazing'

Question 18 *(Prompting and Task Execution)*

You need a 20-page market analysis with competitor profiles, a pricing comparison and recommendations. One big prompt produced a shallow result. What is the best approach?

- **A.** Rewrite the single prompt with every requirement in capitals
- **B.** Break it into steps: agree on an outline, draft each section from its sources, then do a final pass for consistency and tone
- **C.** Regenerate until a deeper version appears
- **D.** Ask for 'much more detail'

Question 19 *(Prompting and Task Execution)*

Claude's first draft of a customer apology is too long and too formal. Which follow-up works best?

- **A.** Regenerate the response
- **B.** Start a new chat with the same prompt
- **C.** 'Cut this to about 120 words, make the tone warmer and more direct, keep the refund details, and apologize in the first sentence.'
- **D.** 'Make it better.'

Question 20 *(Prompting and Task Execution)*

You want campaign ideas for a product launch. Which prompting strategy fits brainstorming?

- **A.** Give no context, so Claude is unconstrained
- **B.** Ask for ideas with final budgets attached
- **C.** Ask for the single best idea
- **D.** Ask for many varied ideas first (say, 20 across different angles), then evaluate and narrow them with your criteria in a follow-up

Question 21 *(Product and Model Selection)*

You need a project timeline that you'll revise with Claude over several rounds and then share with your team. Which output format fits?

- **A.** An Artifact: a standalone deliverable in its own panel that you can refine and share
- **B.** A screenshot of the chat
- **C.** Research mode
- **D.** A plain chat reply you copy out each time

Question 22 *(Product and Model Selection)*

Which TWO tasks justify Opus over Sonnet? (Select TWO)

- **A.** Summarizing a short meeting
- **B.** A strategy recommendation that synthesizes six long reports, where an error would be expensive
- **C.** Writing a two-line reply to a colleague
- **D.** Working out how obligations interact across five related contracts before a high-stakes negotiation
- **E.** Reformatting a list of 30 names into a table

Question 23 *(Product and Model Selection)*

You need to sort 400 short customer comments into five well-defined themes today. Speed matters more than nuance. Which approach fits?

- **A.** Haiku with the five themes clearly defined, followed by a spot-check of a sample
- **B.** One comment per chat, so each gets full attention
- **C.** Opus for maximum quality on every comment
- **D.** Research mode

Question 24 *(Product and Model Selection)*

A long chat about a proposal has become slow, and Claude has started mixing up details from early drafts with the current version. What is the best move?

- **A.** Switch to Haiku for speed
- **B.** Paste the entire chat history into a new chat
- **C.** Start a new chat with a short summary of the current decisions and the latest draft, and keep reference documents in Project knowledge
- **D.** Keep going and correct each mistake as it appears

Question 25 *(Configuration and Knowledge Management)*

In a team Project, where should each of these go: (1) 'Always write in British English and use our headline format'; (2) the 40-page brand guidelines PDF; (3) the brief for this week's campaign?

- **A.** (1) The chat prompt, (2) the chat prompt, (3) Project knowledge
- **B.** (1) Project instructions, (2) the chat prompt, (3) Project knowledge
- **C.** (1) Project knowledge, (2) Project instructions, (3) Project instructions
- **D.** (1) Project instructions, (2) Project knowledge, (3) the chat prompt

Question 26 *(Configuration and Knowledge Management)*

Your weekly report draws on a Google Doc of sales notes that the team updates daily. You've been uploading a copy every Monday, and sometimes it's out of date. What is the better setup?

- **A.** Copy the notes into the Project instructions
- **B.** Connect Google Drive so Claude reads the current version of the document
- **C.** Upload a copy twice a week
- **D.** Paste the notes from memory

Question 27 *(Configuration and Knowledge Management)*

A Project keeps quoting last year's prices. Both the 2025 and the 2026 price sheets are in Project knowledge. What fixes this?

- **A.** Remove or replace the outdated price sheet, and if both must stay, state in the instructions which one is current
- **B.** Use research mode for prices

- **C.** Correct Claude in every chat
- **D.** Switch to Opus

Question 28 *(Troubleshooting and Optimization)*

You asked for 'a table comparing vendors,' and each time the columns come out different. What is the most effective fix?

- **A.** Ask for 'a better table'
- **B.** Show the exact structure: list the columns in order, or paste a sample row, and ask Claude to follow it
- **C.** Regenerate until the columns match
- **D.** Switch to Opus

Question 29 *(Troubleshooting and Optimization)*

Claude answered a question about your company's travel reimbursement limits using general industry norms. The real policy is a PDF on the shared drive. What is the best fix?

- **A.** Tell Claude its answer is wrong and ask it to try again
- **B.** Switch to Opus
- **C.** Attach the policy (or add it to Project knowledge) and ask Claude to answer from it, quoting the relevant section
- **D.** Ask the same question again with more emphasis

Question 30 *(Troubleshooting and Optimization)*

On Monday you changed a Project's instructions, switched the model and uploaded new knowledge files. Since then, the outputs are worse. What should you do?

- **A.** Change everything again
- **B.** Delete the Project and start over
- **C.** Ask Claude what went wrong
- **D.** Roll back, then reapply one change at a time and compare the results on the same test inputs

Answers and explanations

Question 1: A. A grounded claim can point to its evidence, so asking for the clause and reading it yourself is fast and conclusive (A). Asking for reassurance (B) gets a fluent answer either way. Consistency across regenerations (C) doesn't prove accuracy. An attached source (D) lowers the risk of errors without removing it, and this line describes a contract commitment.

Question 2: B. Evaluating output covers bias as well as factual errors, and age-coded language in a job ad is a real fairness and compliance problem (B). D and C change the style and keep the biased phrases. A confuses a common phrase with an acceptable one.

Question 3: C. Precise arithmetic over many figures is a weak fit to accept unchecked, and a client-facing proposal makes the numbers consequential, so they get verified independently (C). The same process checking itself (D) isn't an independent check. Correct inputs (B) don't guarantee correct sums. Rounding (A) hides errors.

Question 4: A and C. A named study with a precise statistic and no source (A), and a confident, specific figure about the future (C), fit the red-flag profile: precision and authority without grounding. A general definition (B) is low risk, and E and D are signs of honest, calibrated output.

Question 5: D. Details beyond the source are a hallucination red flag, and a cancellation claim could trigger real business decisions, so the output has to match its evidence (D). A plausible reading (A) still isn't supported by the transcript. B and C are style issues.

Question 6: B and C. Citations make claims checkable, which helps only if someone checks them (C), and regulatory content goes through a qualified reviewer (B). Citations alone (E) don't prove the summary represents its sources correctly. Repeated claims (D) show consistency, which isn't accuracy. Claude grading its own sources (A) is no independent check.

Question 7: C. A measured pilot shows the real value and the limits on a small scale before the whole company depends on it (C). A big-bang rollout (B) spreads untested habits everywhere at once. D waits for something that won't happen. A produces uneven quality and no evidence.

Question 8: A and E. Drafting variants (E) and repurposing content across channels (A) are classic strong fits, with a person reviewing the output. A spending decision (B) needs human accountability, precise reconciliation at volume (C) belongs in a spreadsheet or finance system, and yesterday's news (D) is beyond the model's knowledge unless research mode grounds it.

Question 9: A. Good stakeholder communication pairs a concrete, measured benefit with an honest limitation (A). B promises something no probabilistic system delivers. C overstates the role and invites resistance. D gives leadership nothing to decide with.

Question 10: D. Solution design starts from the task: what you need, for whom, from which sources, and how you'll judge it (D). Model choice (C) comes later and won't rescue an undefined task. Claude can't infer goals nobody has stated (B). Colors (A) are a detail of the deliverable.

Question 11: B. Claude fits the draft-and-analyze stage, with a qualified person keeping the decision (B). Approval (C), signature (A) and unsupervised negotiation (D) would put legal accountability on a tool.

Question 12: D. Check the policy before pasting: the report is allowed, the customer data isn't, and the task works without it (D). Being internal (B) doesn't change the personal-data rule. Model choice (C) has nothing to do with data handling. Asking Claude to ignore data (A) comes after you've already shared it.

Question 13: A. When the policy doesn't cover a case, escalate to the policy owner and stay on the cautious side in the meantime (A). Silence in the policy (C) doesn't amount to approval. Partial uploads (B) and paraphrasing from memory (D) still share confidential content, and the rewrite adds errors.

Question 14: B. Transparency rules apply as written (B). Editing (D) doesn't remove the requirement, disclosure on request (A) isn't what the policy says, and C misrepresents how the work was made.

Question 15: B and D. Decisions that affect people's livelihoods (B) or legal standing (D) need a qualified human to own them, however good the draft looks. C, A and E are drafting and organizing work where a person still chooses what to use.

Question 16: C. Investment explanations are regulated advice, so qualified review comes before sending (C). A stronger model (A) doesn't change who is accountable. A disclaimer (B) can't replace the review that regulation requires. Speed (D) is beside the point.

Question 17: A and D. Generic output comes from missing context, and the audience and goal (D) and the source material (A) are two of the five parts of a complete business prompt. Enthusiasm (E) adds no information. More versions (C) of an underspecified email are more generic emails. A stronger model (B) still can't know facts you didn't give it.

Question 18: B. Multi-part deliverables work better as a sequence of smaller steps you can check (B). D and A keep the one-giant-prompt problem. C is regenerate-and-hope.

Question 19: C. Specific feedback gives Claude concrete changes to make (C). D gives it nothing to act on, and A and B roll the dice again on an unchanged request.

Question 20: D. Brainstorming works best as volume and variety first, then judgment (D). Asking for one best idea (C) skips the exploration. Budgets (B) belong to a later planning step. With no context (A), the ideas won't fit your product or audience.

Question 21: A. Artifacts are made for standalone deliverables you'll iterate on and share (A). Chat replies (D) suit quick answers. Research mode (C) finds current sources. A screenshot (B) can't be edited.

Question 22: B and D. Opus is for complex, high-stakes reasoning across a lot of material (B, D). Sonnet is the everyday default, and quick, simple tasks like E, C and A suit Sonnet or Haiku.

Question 23: A. Quick, well-defined tasks at volume are Haiku's case, and a sample check keeps the quality honest (A). Opus (C) spends time and cost on a task that doesn't need deep reasoning. Research mode (D) searches for current sources, which this task doesn't need. B wastes hours.

Question 24: C. Restart, summarize and persist is the fix for an overloaded chat: a fresh context with only what matters (C). D keeps the overload, B moves it to a new chat, and A changes speed without fixing the mix-ups.

Question 25: D. Standing rules go in instructions, reference documents go in knowledge, and details for one task go in the chat (D). C swaps rules and documents and makes a one-week brief permanent. A retypes standing rules every time. B re-uploads the guidelines in every chat and stores a temporary brief permanently.

Question 26: B. Connectors suit material that lives in another tool and changes often, so Claude reads the current version (B). More frequent uploads (C) shrink the gap without closing it. Memory (D) adds errors. Instructions (A) hold standing rules, and pasted notes there go stale just as fast.

Question 27: A. Stale knowledge produces confidently outdated answers, so the fix is maintaining the knowledge itself (A). Correcting every chat (C) never ends. A stronger model (D) still reads both sheets. Your prices are internal, so searching for them (B) won't find the right ones.

Question 28: B. A wrong format every time means the format was described vaguely, and a concrete example fixes that (B). C is regenerate-and-hope, D changes the model without clarifying the request, and A is still vague.

Question 29: C. Confident but wrong claims about your own business come from missing source material, so give Claude the policy and ask for the quote (C). Emphasis (D) and a stronger model (B) can't produce a policy Claude hasn't seen. Telling Claude it's wrong (A) without the policy gets you a different guess.

Question 30: D. With three changes at once, you can't tell which one hurt quality, so change one variable at a time against the same inputs (D). A adds more confusion, C asks Claude to explain configuration changes it can't observe, and B throws away a setup that worked before Monday.

Keep going

- **Full-length practice.** The CCAO-F page on Preporato has free sample questions and six full-length timed practice tests (paid) in the exam's format, with a score for every domain.
- **Updates and corrections.** The latest edition, the practice test and the other three guides are at github.com/preporato/claude-certification-guide.